Three improvements to test

Evidence-linked scenarios with explicit elapsed-time assumptions.

Evidence basis: 32 generated requests; 20 in flight; 8 returned at least once. Fixed observation: 23 October 2026, 00:00 UTC. No official pack has been supplied.

Recommendation	Affected	Baseline	Central saving	Sensitivity range
REC-01	4	224 h	112 h	56-145.6 h
REC-02	4	102 h	51 h	25.5-61.2 h
REC-03	4	412 h	103 h	41.2-144.2 h

Read the estimates correctly

Each estimate is the relevant recorded interval multiplied by an assumed reduction. The percentages are planning assumptions, not fitted effects. Open waits stop at the cutoff and are not finished-case durations. The figures are cohort totals, not hours per request. There is no measured HR effort, cost, monthly volume or financial return. Do not add the three headline figures without reconciling overlap.

REC-01 | Catch incomplete submissions before routing

Apply explicit field checks and show a request-type-specific example before submission.

Evidence: 4 affected requests: P2P-1003, P2P-1011, P2P-1019, P2P-1027. 0 h completed-interval time; 224 h open time.

Calculation: 224 relevant elapsed hours × 50% assumed reduction = **112 h**. Sensitivity range: **56-145.6 h**.

Proposed owner: Manager experience and HR policy owner. **Pilot:** Add and version the relevant submission rules. Pilot with comparable request types. Measure return rate and completed turnaround.

REC-02 | Confirm the approval route before sending

Show the proposed human approvers and the organisational evidence used to choose them.

Evidence: 4 affected requests: P2P-1005, P2P-1013, P2P-1021, P2P-1029. 102 h completed-interval time; 0 h open time.

Calculation: 102 relevant elapsed hours × 50% assumed reduction = **51 h**. Sensitivity range: **25.5-61.2 h**.

Proposed owner: HR operations and organisation-design owner. **Pilot:** Display the route and its source. Ask an authorised human to confirm exceptions. Compare routing returns for equivalent requests.

REC-03 | Coordinate unresolved upstream dependencies

Name the upstream request and the person responsible for resolving the blocked transaction.

Evidence: 4 affected requests: P2P-1004, P2P-1012, P2P-1020, P2P-1028. 0 h completed-interval time; 412 h open time.

Calculation: 412 relevant elapsed hours × 25% assumed reduction = **103 h**. Sensitivity range: **41.2-144.2 h**.

Proposed owner: HR operations and upstream transaction owner. **Pilot:** Validate and link the upstream request. Notify both responsible owners. Measure resolved-blocker time in a pilot.

How to validate the scenarios

Replace the demonstration pack with the approved data. Compare equivalent request types in a pilot and record correction returns, completed turnaround, blocked intervals and actual handling effort. Retain the original target calendar and pause rules. Evaluate each intervention before combining its benefits; scale only after monthly volume and adoption are confirmed.

Supporting evidence and reproducibility

The evidence folder includes generated request and event CSVs, complete recommendation JSON, the model report and SHA-256 checksums. Run `node submission/generate-demo-evidence.mjs` from the repository root to reproduce the fixed-cutoff calculations.

Source events by recommendation

REC-01: P2P-1003-E03, P2P-1003-E04, P2P-1011-E03, P2P-1011-E04, P2P-1019-E03, P2P-1019-E04, P2P-1027-E03, P2P-1027-E04

REC-02: P2P-1005-E03, P2P-1005-E04, P2P-1013-E03, P2P-1013-E04, P2P-1021-E03, P2P-1021-E04, P2P-1029-E03, P2P-1029-E04

REC-03: P2P-1004-E04, P2P-1012-E04, P2P-1020-E04, P2P-1028-E04

Official test-set predictions

The organiser historical pack and held-out test set had not been made available to our team at submission time. No official test-set predictions are submitted. The application demonstrates local synthetic scoring with separate delay/return outcomes and evidence. Official prediction evaluation requires the organiser files, approved observation time, mapped schema and historical-data validation.